

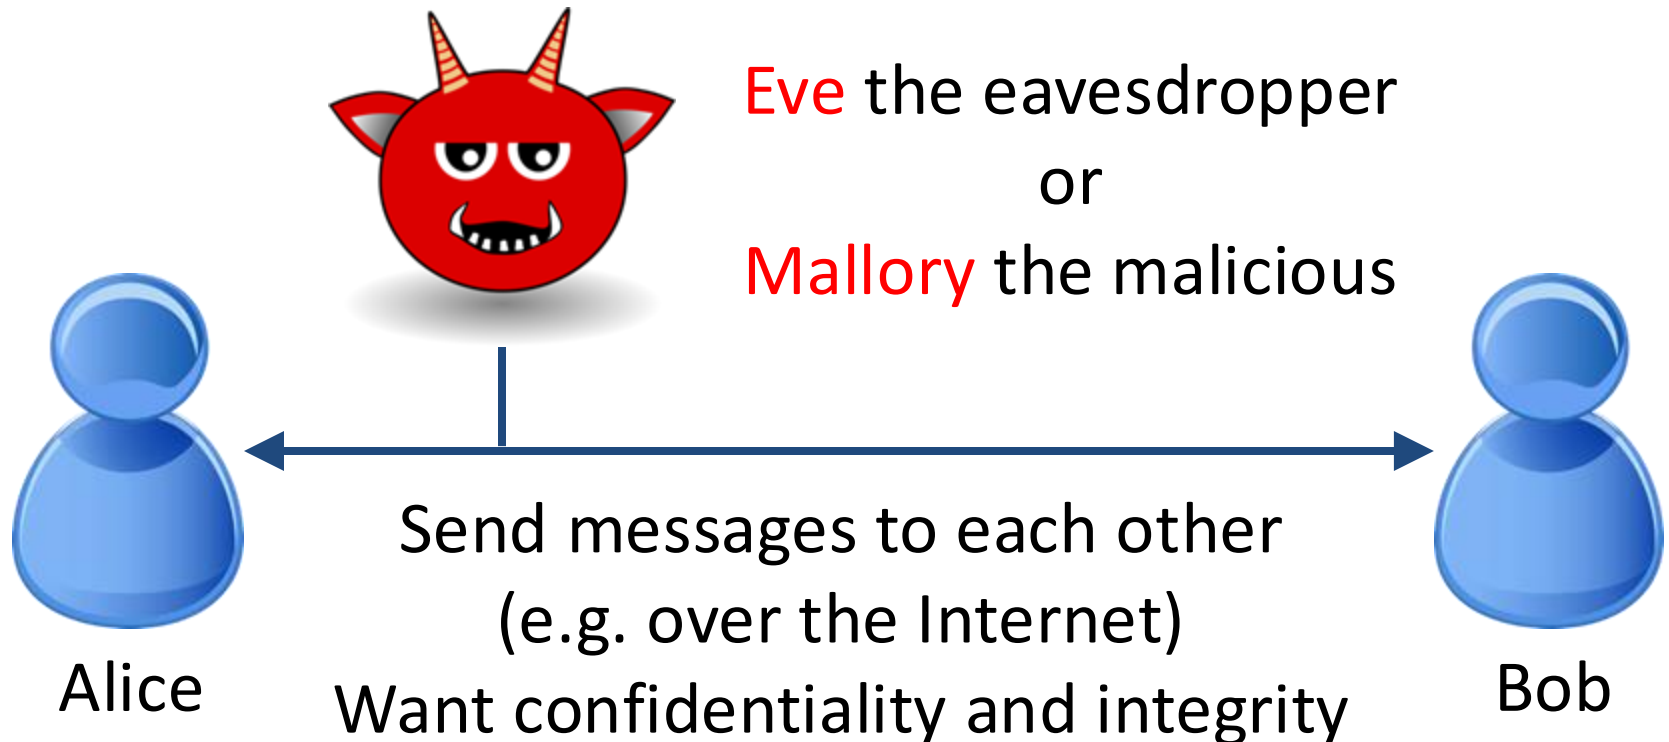
Lecture 15 – Cryptographic Hash Functions

University of Illinois
ECE 422/CS 461

Next 4-5 lectures: Cryptography

Cryptography (or Cryptology)

- Studies techniques for **secure communication** in the presence an **adversary** who has **control** over the **communication channel**



Cryptography (or Cryptology)

- Studies techniques for **secure communication** in the presence an **adversary** who has **control** over the **communication channel**
- Also studies techniques for secure storage, secure collaborative computation, ...

Goals of the Crypto Module

- Primitives we will cover: cryptographic hashing, symmetric & asymmetric encryption, message authentication codes & digital signatures
- Know the **interfaces** of basic crypto primitives
 - What are their inputs and outputs
 - What it means for them to be secure
 - What guarantees they provide and **not** provide
 - Where and how they are typically used
 - Which schemes to use when you need one

Both Rigorous & Empirical

- Modern cryptography is heavily based on *mathematics* but has to resort to *assumptions*
 - Rigorous vs. empirical
- Often *assume* some problem is hard to solve
 - Why do we believe that?
 - Because many experts have tried to solve them for decades or centuries, and could not

Today: Cryptographic Hash Functions

Goals of this Lecture

- By the end of this lecture you should know the following about crypto hash functions:
 - Interface
 - Desired properties
 - Lifecycle and currently recommended ones
 - Common design paradigms
 - Common applications

Cryptographic Hash Functions


- Input – data of an arbitrary length
- Output – fixed length, e.g., 256 bits
- Same input always produces the same output
- Examples: MD5, SHA1, SHA2, SHA3
- $\text{SHA3-256}(\text{"welcome"}) = 64\text{db}51\text{f}8\text{f}79\text{ca}7\text{ec}522\text{a}6\text{b}4\text{ae}5\text{fc}7\text{e}896\text{daac}5318\text{b}2\text{e}82730\text{d}7\text{c}7926\text{b}66\text{d}36\text{eb}$
- $\text{SHA3-256}(\text{"Welcome"}) = 18\text{ec}669\text{de}973\text{b}4483\text{db}9\text{b}64\text{b}2746\text{ceda}564\text{cd}2\text{cdec}2277169382944675\text{a}2\text{ff}9\text{e}$

Applications

- Password hashing
 - System stores (username, salt, $H(\text{pw} || \text{salt})$)
 - User submits (username, pw)
 - System computes $H(\text{pw} || \text{salt})$ and compares

Applications

- Integrity of remote/external storage
 - User computes and stores $H(\text{file})$ locally
 - Compare hash upon download



Desired Properties

- Hard to invert (one-way, OW)
- Hard to find collisions (collision-resistant, CR)
- Exercise: which properties are used in the previous two applications and how?

(Slightly) More Formal Definition

- A cryptographic hash function H with n -bit output is a function:

$$y = H(x): \{0,1\}^* \rightarrow \{0,1\}^n$$

- One-way (OW, also called preimage resistance):
for *almost* all y , **infeasible** to find x s.t. $H(x) = y$
- Collision-resistance (CR): **infeasible** to find x and x' s.t. $x \neq x'$ and $H(x) = H(x')$

What Does “Infeasible” Mean?

- Infeasible \neq impossible
- In fact, both inversion and collision-finding are clearly possible by brute-force
 - Collisions must exist due to pigeon-hole principle
 - Brute-force collision in $O(2^{n/2})$ time and space due to “birthday paradox”
 - Brute-force inversion in $O(2^n)$ time
- Infeasible = no known attacks (yet) better than brute-force attacks

How to Choose n ?

- Make $2^{n/2}$ a prohibitive cost
 - Since collision is the easier brute-force attack
- $n = 128$ used to be popular but is now too small
 - 2^{64} is no longer prohibitive
- Typical choice: $n = 160, 192, 224$ or 256
- $n = 384$ or 512 for the paranoid
 - $2^{256} \approx$ the number of particles in the universe

Desired Properties

- Do OW and CR imply each other?
 - No. We can construct hash functions that have one property but not the other. Fun challenge :)
- A stronger property is **pseudorandom**: for every input x , $H(x)$ “looks random”
 - Implies OW and CR
 - But at the same needs to be deterministic ...

(Slightly) More Formal Definition


- Ideal and unattainable hash: **random oracle**
 - Maintain a (infinite sized) table for every input-output pairs (x, y)
 - On new input x , generate a random n -bit value y , store (x, y) in table, and output y
 - On previous input x , output the stored y
- H is pseudorandom if H “behaves like” a random oracle

Well-known Crypto Hash Functions

- ~~MD2, MD4, MD5, SHA1~~, SHA2, SHA3, ...
 - ~~Strikethrough~~ = broken, never use again!
- How are they designed?
- How do we know they are OW and CR?
- Experts use their insights and experience to design and inspect/attack each other's design
 - MD = Message Digest, a series by Ron Rivest
 - SHA = Secure Hash Algorithm, NIST competition

Do NOT roll your own crypto!

Lifecycle of Crypto Hash Functions



[1] Note that 128-bit hashes are at best 2^{64} complexity to break; using a 128-bit hash is irresponsible based on sheer digest length.

[2] What happened in 2004? [Xiaoyun Wang and Dengguo Feng and Xuejia Lai and Hongbo Yu](#) happened.

[3] Google spent [6500 CPU years and 110 GPU years](#) to convince everyone we need to stop using SHA-1 for security critical applications. Also because it was cool.

[4] In 2007, the [NIST launched the SHA-3 competition](#) because "Although there is no specific reason to believe that a practical attack on any of the SHA-2 family of hash functions is imminent, a successful collision attack on an algorithm in the SHA-2 family could have catastrophic effects for digital signatures." One year later the [first strength reduction](#) was published.

- Eventually a function weakened
- Time to move to a new function and (hopefully) stay ahead of attackers (before a collision is found)

Common Design Paradigm

- How do we design a function that takes arbitrarily long input?
- Merkle-Damgård construction
 - Adopted by ~~MD5, SHA1~~, SHA2
 - First, design a compressing function that takes fixed-length inputs


$$\boxed{f} : \{0,1\}^{2n} \rightarrow \{0,1\}^n$$

- Then, “absorb” the input n-bit at a time

SHA2-256 Compressing Function

- Intentionally “hairy” and “messy”
- 64 rounds of this

$$\begin{aligned}\text{Ch}(E, F, G) &= (E \wedge F) \oplus (\neg E \wedge G) \\ \text{Ma}(A, B, C) &= (A \wedge B) \oplus (A \wedge C) \oplus (B \wedge C) \\ \Sigma_0(A) &= (A \ggg 2) \oplus (A \ggg 13) \oplus (A \ggg 22) \\ \Sigma_1(E) &= (E \ggg 6) \oplus (E \ggg 11) \oplus (E \ggg 25)\end{aligned}$$



$$\{0,1\}^{512} \rightarrow \{0,1\}^{256}$$

Merkle-Damgård Construction

Arbitrary
length input x




$x || \text{padding}$
($||$ means bit concat)

Partition into
 n -bit chunks



$\{0,1\}^{2n} \rightarrow \{0,1\}^n$

IV: some fixed
 n -bit value




Final n -bit output: $H(x)$

Merkle-Damgård Construction

- Theorem: if the compressing function f is one-way (OW) and collision-resistant (CR), then a Merkle-Damgård hash is also OW and CR
- Proof idea: suppose some attacker breaks Merkle-Damgård, it also breaks f
- This is called security reduction. Allows us to focus on security of basic building blocks.

Reduction Proofs


- f OW \rightarrow Merkle-Damgård OW
 - Suppose Merkle-Damgård is not OW, i.e., there is a feasible algo that finds preimage x s.t. $H(x) = y$
 - Easy to evaluate f “forward”
 - Then, $f(u || b_{m-1}) = y$. Preimage found for y under f . QED



Reduction Proofs

- f OW \rightarrow Merkle-Damgård OW

- “ f CR \rightarrow Merkle-Damgård CR” is also true with “proper” padding




Subtleties in Padding

- Recall that we pad input to a multiple of n bits
- Pad with 0?
 - SomeLongBitString0000000
 - SomeLongBitString0000000
- Merkle-Damgård requires that inputs with different lengths, once padded, differ in their last blocks (typically, embed length in padding)
 - Can now prove CR reduction (exercise)

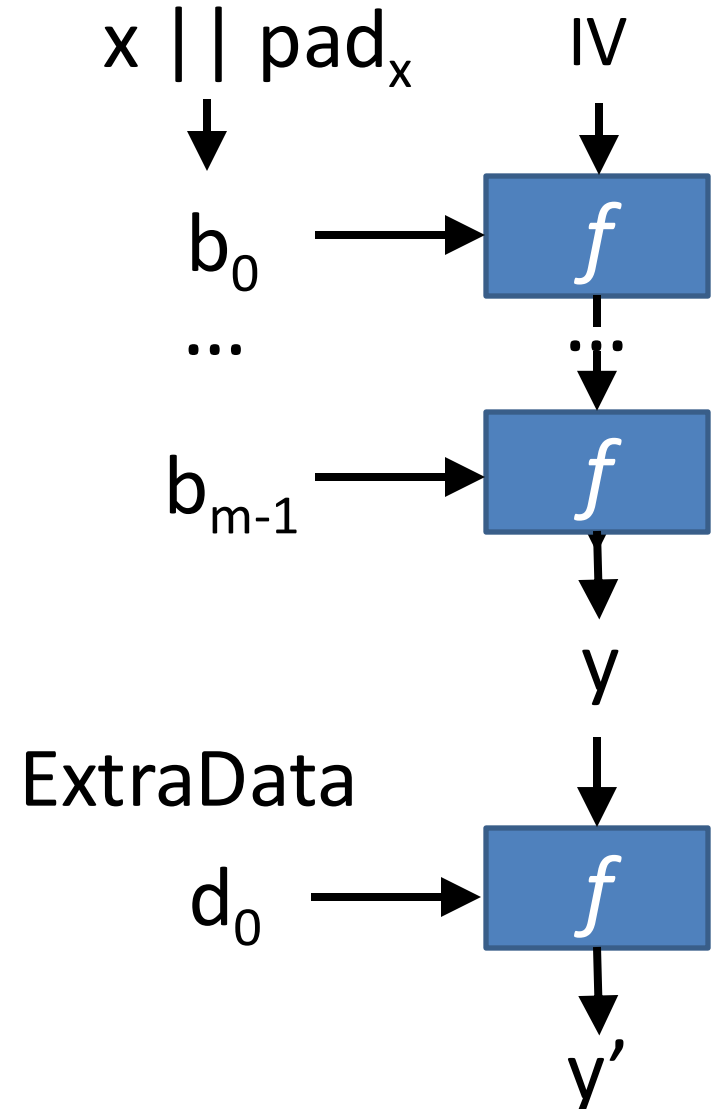
Merkle-Damgård Construction

- f OW \rightarrow Merkle-Damgård OW
- f CR \rightarrow Merkle-Damgård CR
(with proper padding)
- How about pseudorandomness?
 - Merkle-Damgård is NOT pseudorandom even if f is



Length Extension Attack

- Given $H(x)$, one can compute $H(x || \text{pad}_x || \text{ExtraData})$ with more rounds of f
- A random oracle would not exhibit this behavior
 - Given $\{ H(x_i)=y_i \}$, for a new x' , $H(x')$ would be random




Length Extension Attack

- Applicable to all Merkle-Damgård constructions including ~~MD5, SHA1~~, SHA2
- Not a show-stopper as they do not affect OW and CR
- If you need pseudorandomness, use SHA3 (not a Merkle-Damgård construction)


SHA3: Sponge Construction

- Final hash output does not expose the entire internal state \rightarrow cannot length-extend
- Along with other techniques, achieves pseudo-randomness (if f is pseudorandom)

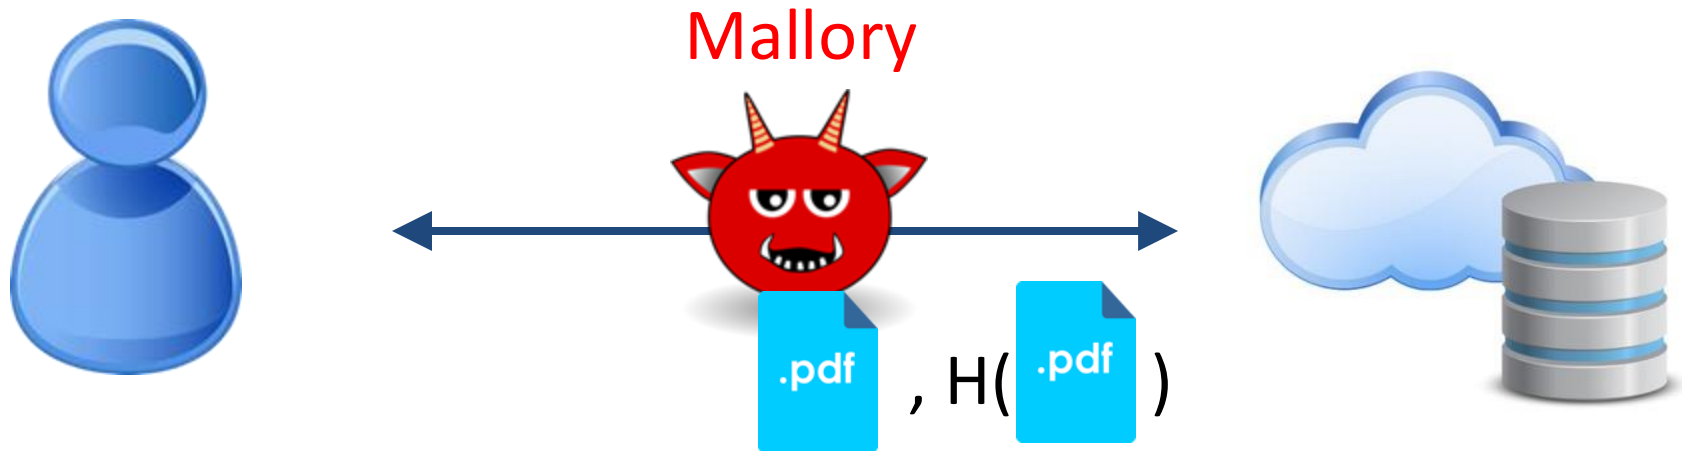


Recall External Storage Application

- Integrity of remote/external storage
 - User computes and stores $H(\text{file})$ locally
 - Compare hash upon download
 - Collision-resistance protects integrity

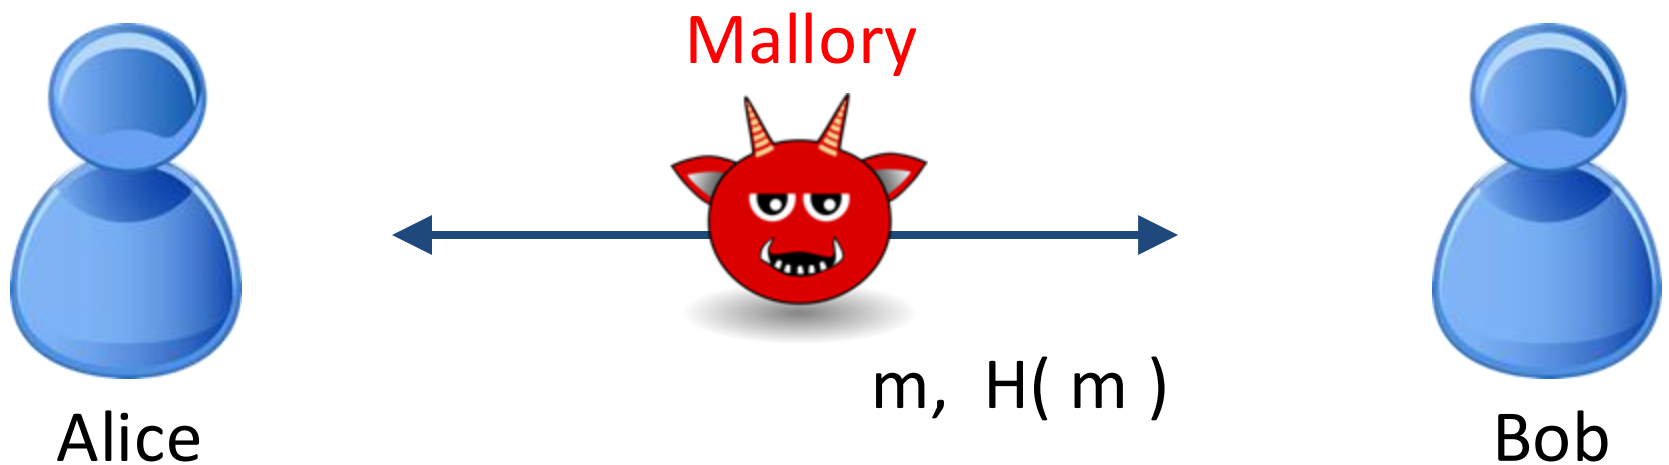


Integrity of Download?



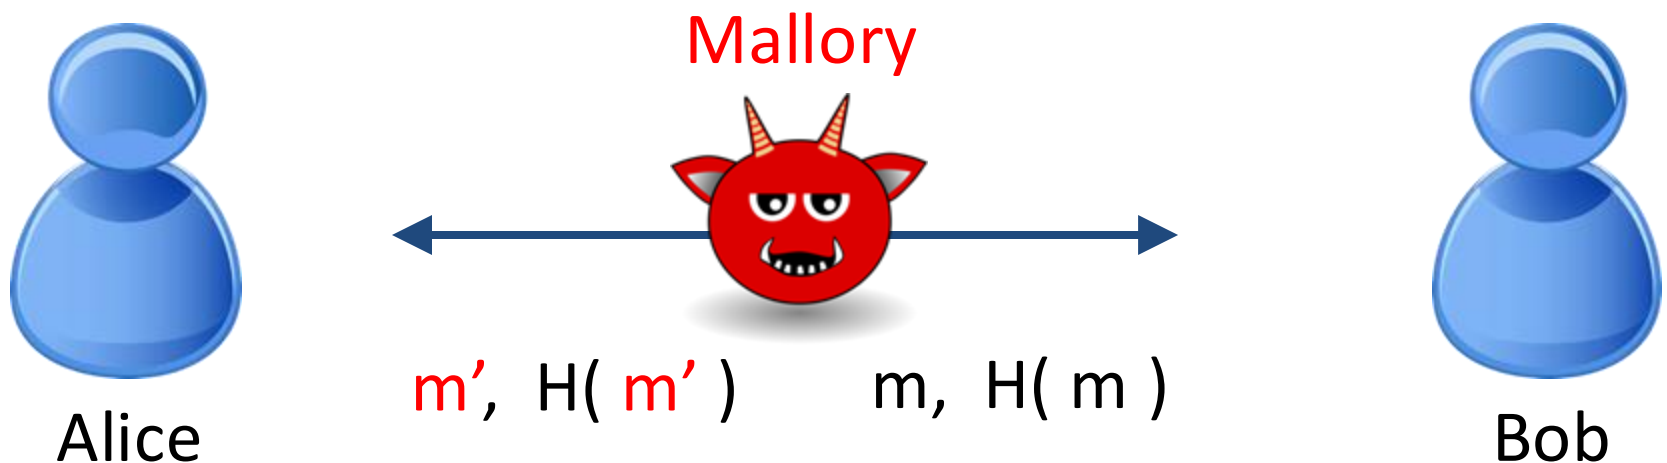
Integrity of Communication?

- I.e., message authentication




Message Authentication

- In its simplest form, a cryptographic hash does NOT work for a message authentication. Attacker can hash the modified message.



Message Authentication

- Two settings it can work:
 - If the hash can be transmitted in another *trustworthy* but low-bandwidth channel



Message Authentication

- Two settings it can work:
 - If the hash can be transmitted in another *trustworthy* but low-bandwidth channel, or
 - If Alice and Bob share a secret key k
 - Will come back to this in a future lecture



Alice
(knows k)



Bob
(knows k)

Another Application

- Let's play a game online: if you guess my favorite 2-digit number in one try, you get A+
 - You will never win 😊
- Need a “sealed envelop” for a fair game
 - Alice sends Bob $c = \text{Commit}(m)$
 - Alice later “opens” m for Bob to verify against c
 - Hiding: Bob cannot find out m from c
 - Binding: Alice cannot open to another $m' \neq m$

Commitment

- Alice sends Bob $c = \text{Commit}(m) = H(m || r)$
where r is a long, fixed-length & random string
 - Why do we need r ?
- Alice can later reveal m and r to “open”
- Hiding: Bob cannot find m from c
 - If H is pseudorandom (OW is insufficient, why?)
- Binding: Alice cannot open to another $m' \neq m$
 - If H is collision-resistant

Summary

- Cryptographic hash function:
 - Definition $H: \{0,1\}^* \rightarrow \{0,1\}^n$
 - Desired properties: one-way, collision-resistant, pseudorandom (behave like a random oracle)
 - Currently recommended: SHA3 (SHA2 if must)
 - Paradigms: Merkle-Damgård and sponge
 - Note length extension attacks for Merkle-Damgård!
 - Applications: password hashing, external storage, commitment, hash-based message authentication (need extra assumptions), ...